

evidence-gamma-normal

April 7, 2022

```
[1]: import torch
import numpy as np
import matplotlib.pyplot as plt

plt.rc('text', usetex=True)
plt.rc('font', **{'family': 'sans-serif', 'serif': ['Palatino']})
figSize = (12, 8)
fontSize = 20
```

1 Notations

Similar to [wikipedia](#).

$$X \sim \Gamma(\alpha, \beta)$$

The probability distribution function is given by

$$f(x; \alpha, \beta) = \frac{x^{\alpha-1} e^{-\beta x} \beta^\alpha}{\Gamma(\alpha)}$$

where $\Gamma(\alpha) = (\alpha - 1)!$

α is the shape parameter.

β is the rate parameter (we will fix this parameter to $\beta = 5.0$ throughout this work).

2 PyTorch Notations

[Link](#)

```
torch.distributions.gamma.Gamma(concentration, rate, validate_args=None)
```

Creates a Gamma distribution parameterized by shape concentration and rate.

3 Analysis

3.0.1 Steps

- The data is generated from the Gamma distribution.

- The parameter β is fixed to $\beta = 5.0$.
- Only one parameter, α varied.
- Assume a normal prior on α , that is, $p(\alpha) = \mathcal{N}(\alpha_*, 1)$
- Goal: compute evidence with and without compression.
- Goal: use the score function for compression.
- Goal: compare with the case when a Gaussian likelihood is used.

3.0.2 Bayes' theorem

$$p(\alpha, \beta | \mathbf{x}) = \frac{p(\mathbf{x} | \alpha, \beta) p(\alpha, \beta)}{p(\mathbf{x})}$$

$p(\beta) = \delta(\beta - 5.0)$. We will ignore β below since it is fixed.

$$p(\alpha | \mathbf{x}) = \frac{p(\mathbf{x} | \alpha) p(\alpha)}{p(\mathbf{x})}$$

Without Compression

$$p(\mathbf{x}) = \int \Gamma(\alpha, \beta = 5.0) p(\alpha) d\alpha$$

Assume we have generated N iid samples, $\mathbf{x} \in \mathbb{R}^N$:

$$p(\mathbf{x}) = \int \prod_{i=1}^N f(x_i; \alpha, \beta = 5.0) p(\alpha) d\alpha$$

We can do this integration numerically (on a grid).

If we assume a Gaussian likelihood,

$$p(\mathbf{x}) = \int \prod_{i=1}^N f(x_i; \mu, \sigma = 1.0) p(\mu) d\mu$$

With Compression We will be using the score function, that is,

$$s = \nabla \mathcal{L}$$

where \mathcal{L} is the log-likelihood. We need the derivative of the log-likelihood with respect to α , but evaluated at each sample. The log-likelihood when using the Gamma distribution is:

$$\mathcal{L} = \sum_{i=1}^N \log f(x_i; \alpha, \beta = 5.0)$$

and the derivative of the log-likelihood with respect to the parameter, α is:

$$\nabla \mathcal{L} = \nabla \sum_{i=1}^N \log f(x_i; \alpha, \beta = 5.0)$$

Therefore, we have compressed N samples (data points) to just 1 number. The likelihood (for the compressed data) is

$$p(s|\alpha) = \frac{1}{\sqrt{2\pi\sigma_*^2}} \exp \left[-\frac{1}{2} \frac{(s - s_*)^2}{\sigma_*^2} \right]$$

where $s_* = \nabla \mathcal{L}(\alpha = \alpha_*)$ and $\sigma_*^2 = -\nabla^2 \mathcal{L}(\alpha = \alpha_*)$. α_* is a chosen fiducial point.

3.0.3 Explore the Gamma Distribution


```
[2]: shape = 10.0

# the rate is fixed
rate = 5.0

dist = torch.distributions.gamma.Gamma(shape, rate)

samples = dist.rsample([10000]).view(-1)

[3]: plt.figure(figsize=(8,8))
plt.hist(samples, density=True, bins=20, ec='b', alpha=0.5, color='skyblue')
plt.ylabel(r'$p(x)$', fontsize = fontSize)
plt.xlabel(r'$x$', fontsize = fontSize)
plt.tick_params(axis='x', labelsize=fontSize)
plt.tick_params(axis='y', labelsize=fontSize)
plt.show()
```



3.0.4 Data

This is our data vector of size 500 and is fixed throughout the analysis.

```
[4]: data = dist.rsampl([300]).view(-1)
```

4 Gamma Distribution Likelihood

4.0.1 Parameter Estimation Case

We can infer α from the data (sanity check).

```
[5]: def loglike_gamma(data, alpha: float, rate: float = 5.0):
```

```

alpha = torch.tensor([alpha], dtype = torch.float32)

rate = torch.tensor([rate], dtype = torch.float32)

dist = torch.distributions.gamma.Gamma(alpha, rate)

logp = dist.log_prob(data)

loglike = torch.sum(logp).item()

return loglike

```

```
[6]: def logprior_normal(alpha, mean: float = 10.0, sigma: float = 1.0):
```

```

    alpha = torch.tensor([alpha], dtype = torch.float32)

    mean = torch.tensor([mean], dtype = torch.float32)

    sigma = torch.tensor([sigma], dtype = torch.float32)

    dist = torch.distributions.normal.Normal(mean, sigma)

    logprior = dist.log_prob(alpha).item()

    return logprior

```

```
[7]: npoint = 1000
alpha_grid = torch.linspace(9.0, 11.0, npoint)
logpost_gamma = torch.zeros(npoint)
```

```
[8]: for i in range(npoint):
```

```

    # the value of alpha
    alpha_value = alpha_grid[i].item()


    # the log-posterior value
    logpost_gamma[i] = loglike_gamma(data, alpha_value)
    logpost_gamma[i] += logprior_normal(alpha_value)

```

```
[9]: # the pdf normalised by the maximum
pdf_gamma = torch.exp(logpost_gamma - logpost_gamma.max())
```

```
[10]: plt.figure(figsize = (8,8))
plt.plot(alpha_grid, pdf_gamma, lw = 2.0)
plt.ylabel(r'$p(\alpha)$', fontsize = fontSize)
plt.xlabel(r'$\alpha$', fontsize = fontSize)
```

```
plt.tick_params(axis='x', labelsz=fontSize)
plt.tick_params(axis='y', labelsz=fontSize)
plt.ylim(0.0, 1.2)
plt.xlim(min(alpha_grid), max(alpha_grid))
plt.show()
```



4.0.2 Summary Statistics

```
[11]: pdf_gamma_norm = pdf_gamma/ torch.trapz(pdf_gamma, alpha_grid)
```

We can calculate the mean and variance as follows:

$$\mathbb{E}[\alpha] = \int \alpha p(\alpha|\mathbf{x}) d\alpha$$

$$\text{var}[\alpha] = \int \alpha^2 p(\alpha|\mathbf{x}) d\alpha - \mathbb{E}^2[\alpha]$$

```
[12]: mean_alpha = torch.trapz(alpha_grid*pdf_gamma_norm, alpha_grid)
```

```
[13]: var_alpha = torch.trapz(alpha_grid**2 * pdf_gamma_norm, alpha_grid) -
      ↪ mean_alpha**2
```

```
[14]: print(f'The mean is {mean_alpha.item():.4f}')
      print(f'The standard deviation is {var_alpha.sqrt().item():.4f}')
```

The mean is 9.8626

The standard deviation is 0.1741

Good - we expected a value around 10 since the samples (data) were generated from $X \sim \Gamma(10.0, 5.0)$.

4.0.3 Evidence Estimate

We basically need to do a 1D-integration of the posterior (up to a normalisation constant).

```
[15]: evi_gamma = torch.trapz(pdf_gamma, alpha_grid)
```

```
[16]: log_evi_gamma = logpost_gamma.max().item() + torch.log(evi_gamma)
```

```
[17]: log_evi_gamma
```

```
[17]: tensor(-283.0836)
```

5 Normal Distribution Likelihood

5.0.1 Parameter Estimation Case

```
[18]: def loglike_normal(data, mean, sigma: float = 1.0):

    mean = torch.tensor([mean], dtype = torch.float32)

    sigma = torch.tensor([sigma], dtype = torch.float32)

    dist = torch.distributions.normal.Normal(mean, sigma)

    logp = dist.log_prob(data)

    loglike = torch.sum(logp).item()

    return loglike
```

Looking at the Gamma distribution above, the mean is roughly centered on 2.

```
[19]: npoint = 1000
mu_grid = torch.linspace(1.5, 2.5, npoint)
logpost_normal = torch.zeros(npoint)
```

```
[20]: for i in range(npoint):


    # the value of alpha
    mu_value = mu_grid[i].item()

    # the log-posterior value
    logpost_normal[i] = loglike_normal(data, mu_value)

    # add the prior contribution with a mean = 2.0 and standard deviation = 1.0
    logpost_normal[i] += logprior_normal(mu_value, 2.0, 1.0)
```

```
[21]: # the pdf normalised by the maximum
pdf_normal = torch.exp(logpost_normal - logpost_normal.max())
```

```
[22]: plt.figure(figsize = (8,8))
plt.plot(mu_grid, pdf_normal, lw = 2.0)
plt.ylabel(r'$p(\mu)$', fontsize = fontSize)
plt.xlabel(r'$\mu$', fontsize = fontSize)
plt.tick_params(axis='x', labelsize=fontSize)
plt.tick_params(axis='y', labelsize=fontSize)
plt.ylim(0.0, 1.2)
plt.xlim(min(mu_grid), max(mu_grid))
plt.show()
```

```
[23]: evi_normal = torch.trapz(pdf_normal, mu_grid)
```

```
[24]: log_evi_normal = logpost_normal.max().item() + torch.log(evi_normal)
```

```
[25]: log_evi_normal
```

```
[25]: tensor(-339.0552)
```

```
[26]: logBF = log_evi_gamma - log_evi_normal
      print(f'The log-Bayes factor between Model 1 (Gamma) and Model 2 (Normal) is_
            ↪{logBF:.2f}')
```

The log-Bayes factor between Model 1 (Gamma) and Model 2 (Normal) is 55.97

Good - we expect the Gamma likelihood function to be better compared to a Gaussian likelihood

function.

6 Compression - Score Function

```
[27]: def score_compression_gamma(data, alpha: float, rate: float = 5.0):  
  
    alpha = torch.tensor([alpha], dtype = torch.float32)  
  
    # we need the derivatives with respect to the shape parameter  
    alpha.requires_grad = True  
  
    rate = torch.tensor([rate], dtype = torch.float32)  
  
    dist = torch.distributions.gamma.Gamma(alpha, rate)  
  
    # the log-likelihood  
    log_like = dist.log_prob(data).sum()  
  
    # gradient of log-likelihood with respect to shape (alpha)  
    grad_log_like = torch.autograd.grad(log_like, alpha, create_graph=True)[0]  
  
    # the second derivative with respect to shape (alpha)  
    hessian = torch.autograd.grad(grad_log_like, alpha)[0]  
  
    return grad_log_like.item(), hessian.item()
```

We will choose $\alpha = 10.0$ as the expansion (fiducial) point. We can also quickly compute the Maximum Likelihood Estimator (MLE), that is, find α_{MLE} for which $\nabla \mathcal{L} = 0$ and this happens at:

$$\alpha_{\text{MLE}} = \frac{\beta}{N} \sum_{i=1}^N x_i$$

In our work, $\beta = 5.0$. So we want just compute α_{MLE} .

```
[28]: # Maximum likelihood estimate  
torch.mean(data).item() * 5.0
```

```
[28]: 9.87557053565979
```

```
[29]: # compress the data at the shape parameter = 10.0  
  
s_data, h_data = score_compression_gamma(data, 10.0)
```

```
[30]: print(f'Error estimate of the score: {np.sqrt(-1.0 / h_data):.4f}')
```

Error estimate of the score: 0.1780

```
[31]: print(f'Previously, we estimated the standard deviation of the shape parameter_
      ↪using the posterior as: {var_alpha.sqrt().item():.4f}')
```

Previously, we estimated the standard deviation of the shape parameter using the posterior as: 0.1741

7 ISSUE below

The pdf is *always* centred on the expansion point? and the error estimate on α does not look good!

```
[32]: def loglike_gamma_score(alpha, data, score, hessian, rate: float = 5.0):

      grad, _ = score_compression_gamma(data, alpha, rate)

      # print(grad)

      # calculate the variance (inverse of hessian)
      var = torch.tensor([-1. / hessian])

      std = torch.sqrt(var)

      gaussian = torch.distributions.normal.Normal(score, std)

      loglike = gaussian.log_prob(torch.tensor([grad]))

      return loglike
```

```
[33]: logpost_gamma_score = torch.zeros(npoint)
```

```
[34]: alpha_test = torch.linspace(9.90, 10.10, 1000)
      for i in range(npoint):


          # the value of alpha
          alpha_value = alpha_test[i].item()

          # the log-posterior value
          logpost_gamma_score[i] = loglike_gamma_score(alpha_value, data, s_data,
      ↪h_data, rate = 5.0)
          logpost_gamma_score[i] += logprior_normal(alpha_value, 10.0, 1.0)
```

```
[35]: # the pdf normalised by the maximum
      pdf_gamma_score = torch.exp(logpost_gamma_score - logpost_gamma_score.max())
```

```
[36]: plt.figure(figsize = (8,8))
      plt.plot(alpha_test, pdf_gamma_score, lw = 2.0)
      plt.ylabel(r'$p(\alpha)$', fontsize = fontSize)
      plt.xlabel(r'$\alpha$', fontsize = fontSize)
```

```
plt.tick_params(axis='x', labelsz=fontSize)
plt.tick_params(axis='y', labelsz=fontSize)
plt.ylim(0.0, 1.2)
# plt.xlim(min(alpha_grid), max(alpha_grid))
plt.show()
```




7.0.1 Test

We can just sample the score, given the score function calculated at $\alpha = 10.0$ and the error estimate. This looks good but we would expect the error estimate on α to be of the same width (since the Hessian is just the second derivative of the log-likelihood with respect to the parameter α). Maybe not - maybe we need a lot of data for the likelihood to peak. Maybe the likelihood surface is too broad.

```
[37]: gaussian_test = torch.distributions.normal.Normal(s_data, np.sqrt(-1.0 /  $\lambda$ 
↪h_data))
```

```
[38]: plt.figure(figsize=(8,8))
plt.hist(gaussian_test.rsample([10000]), density=True, bins=20, ec='b', alpha=0.
↪5, color='skyblue')
plt.ylabel(r'$p(s)$', fontsize = fontSize)
plt.xlabel(r'$s$', fontsize = fontSize)
plt.tick_params(axis='x', labels=fontSize)
plt.tick_params(axis='y', labels=fontSize)
plt.show()
```



```
[ ]:
```